

# COMS E6998: Microservices and Cloud Applications

*Lecture 1: Introduction*

Dr. Donald F. Ferguson  
[dff9@columbia.edu](mailto:dff9@columbia.edu)

# A Little About Me

 The image part with relationship ID rId2 was not found in the file.

 The image part with relationship ID rId2 was not found in the file.

# Contents

# Contents

- Introduction
  - A bit about your instructor.
  - Coursework and grading.
  - Logistics.
- Course Overview and Contents
  - Schedule and topics.
  - Cloud
  - Microservices
  - Cloud applications.
- An example “microservice, cloud application”
  - Demo.
  - Overall structure.
- 1<sup>st</sup> Assignment

# Introduction

# Your Instructor

# A Little About Me

- Columbia through and through
  - BA (Columbia College) '82, M.S. '84 (Engineering), M.Phil. '84 (GSAS).
  - Ph.D. '87 (GSAS) – Thesis: "The Application of Microeconomics to the Design of Resource Allocation and Control Algorithms."
  - Adjunct Professor at Columbia Univ: 2006, 2009, 2014-2017. 7 "Topics in CS" courses.
- 32 years of technology contribution and leadership.
  - VP, IBM Fellow: IBM Research, IBM Software Group.
  - Microsoft Technical Fellow.
  - Distinguished Engineer, CTO, Executive VP, CA technologies.
  - VP, Sr. Fellow, CTO, Dell Software Group.
  - Co-Founder, CTO Seeka TV ([www.seekatv.com](http://www.seekatv.com)).
- Personal
  - Two fantastic daughters; high school and one is a senior at Barnard.
  - Black belt in Kenpo Karate (<http://www.tracyskarateworldwide.com/about.html>, [www.elitedefensivetactics.com](http://www.elitedefensivetactics.com))
  - 2LT in New York Guard (<http://dmna.ny.gov/nyg/>)
  - Conversational Spanish, learning Arabic
  - Hobbies: Bicycling, Astronomy

# A Little About Me

- Columbia through and through
  - BA (Columbia College) '82, M.S. '84 (Engineering), M.Phil. '84 (GSAS).
  - Ph.D. '87 (GSAS) – Thesis: "The Application of Microeconomics to the Design of Resource Allocation and Control Algorithms."
- 32 years of technology contribution and leadership.
  - IBM Fellow: IBM Research, IBM Software Group.
  - Microsoft Technical Fellow.
  - Distinguished Engineer, CTO, Executive VP, CA technologies.
  - Sr. Fellow, CTO, Dell Software Group.
  - Co-Found, CTO Seeka TV.
- Personal
  - Two daughters
  - Black belt in Kenpo Karate
  - 2LT in New York Guard (<http://dmna.ny.gov/nyg/>)
  - Conversation Spanish, learning Arabic
  - Hobbies: Bicycling, Astronomy

# A Lot of Experience

 The image part with relationship ID rId2 was not found in the file.

# A Lot of Experience

 The image part with relationship ID rId2 was not found in the file.

# A Lot of Experience

I may not be able to teach you the right  
way to do something,  
but  
I can show you plenty of  
wrong ways to avoid.

# Coursework Grading

# Coursework and Grading

- No midterm, final or exams.
- Five 20 point projects, each of which extends previous projects.
  - Approximately every two weeks, with two weeks to complete.
  - Teams
    - Approximately 5 team members.
    - Each team has a designated *primary contact* and *team name*.
  - Each project has the following sub-deliverables, packaged in a zip file with a readme.
    - A set of user stories.
    - A short, simple architecture/design description with diagrams.
    - Demo
    - Application artefacts, e.g. code, DB schema/data, HTML, URL for testing, ...
  - There will be a simulated “project review,” which is a meeting with presentation, demo, Q&A, etc.
- In addition to the technology, I am trying to *mentor* you to learn how to be a successful technical professional. You will experience very basic concepts in
  - Product management.
  - Presentation and documentation skills.
  - Transforming vague requirements into something clear and implementable.
  - Q&A and discussion of alternatives.
  - Teamwork and team dynamics.

This will drive you  
crazy in the  
beginning.

# Logistics

# Logistics

- Course material
  - *No textbooks.* Technology changes too quickly, and we will cover concepts and technologies from several disciplines.
  - Lecture and lecture notes will provide links to online information.
  - Web search is your first and best resource.
- CAs:
  - We have two excellent CAs, with whom I have previous worked.
  - Will introduce themselves via email, Piazza and other collaboration environments.
- Office hours
  - Professor: Tu/Th 10:15 to 12:00, or by appt (typically Google Hangout).
  - CAs: TBA
- Collaboration environments:
  - CourseWorks for submitting projects.
  - Piazza: Slides, discussion ([piazza.com/columbia/fall2017/comse6998\\_014\\_2017\\_3topicsincomputerscience](https://piazza.com/columbia/fall2017/comse6998_014_2017_3topicsincomputerscience))
  - Slack: Good way to have short discussions, get ahold of me, instant messages, etc.
    - Team: ([https://join.slack.com/t/dff-columbia/shared\\_invite/MjM1MjA5NjAwOTk2LTE1MDQyODExOTQtNzEzM2FjZmMzMw](https://join.slack.com/t/dff-columbia/shared_invite/MjM1MjA5NjAwOTk2LTE1MDQyODExOTQtNzEzM2FjZmMzMw)).
    - Join channel #e6998-2017
- We will use Amazon Web Services, but the concepts are general.
- **Questions and discussion are expected, appreciated and essential.**

# Course Overview and Content

# Course Overview

- Course Overview
- Core Concepts

- Fundamentals
- REST
- Multi-Tenant

- Middleware
- Session

- Oauth 2.0
- Basic Security

- Calling Cloud APIs
- Integration Patterns

- Pub/Sub
- Message Queues

- Serverless
- API Gateway

- Dynamo DB
- Neo4J
- Redis

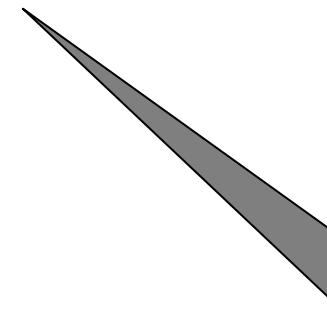
- Data Pipeline
- Step Functions

School holiday.

- Web Hooks
- XMPP/Chat

- Rules Engine

- Text
- Search

- 
- **Overflow**
  - **Discussion**

# Comments

- This course is about building applications
    - That apply best practices and design patterns,
    - Using a broad set of technologies,
    - Each of which is several lectures or a complete course unto itself.  
We will cover the essentials but not the depth of the individual technologies.
  - One place where we will "go deep" is on the basic structure of a well-designed application/set of microservices. Some example topics are
    - Model-View-Controller.
    - Application Tier-Data Tier.
    - Encapsulation/Abstraction.
    - REST API
- My experience from prior courses has been that I have gone through these topics too quickly, causing students to do advanced topics with a poor design.*

# Course Content:

Cloud  
Microservices  
Applications

# Cloud

# IaaS, PaaS, SaaS

 The image part with relationship ID rld2 was not found in the file.

# IaaS, PaaS, SaaS

 The image part with relationship ID rld2 was not found in the file.

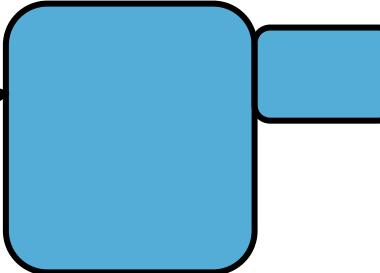
Focus will be on building

- Microservices.
- Cloud applications.

Using

- Application containers.
- Middleware and MW APIs.
- SaaS APIs

And less on building containers  
(e.g. Docker) and lower layer  
infrastructure (e.g. Kubernetes)



# IaaS, PaaS, SaaS – A Little More about PaaS

# IaaS, PaaS, SaaS – A Little More about PaaS

The image part with relationship ID rId2 was not found in the file.

Additionally

- Pub/Sub
- Several data service models
- Web callable services, e.g. FB
- Rules engine
- Design patterns
- ... ... TBD.

# Cloud is about \*- as-a-Service

([https://en.wikipedia.org/wiki/As\\_a\\_service](https://en.wikipedia.org/wiki/As_a_service))

“\*aaS is an acronym for **as a service**, and refers to something being made available over the Internet to a customer as a service.<sup>[1]</sup> Examples include:

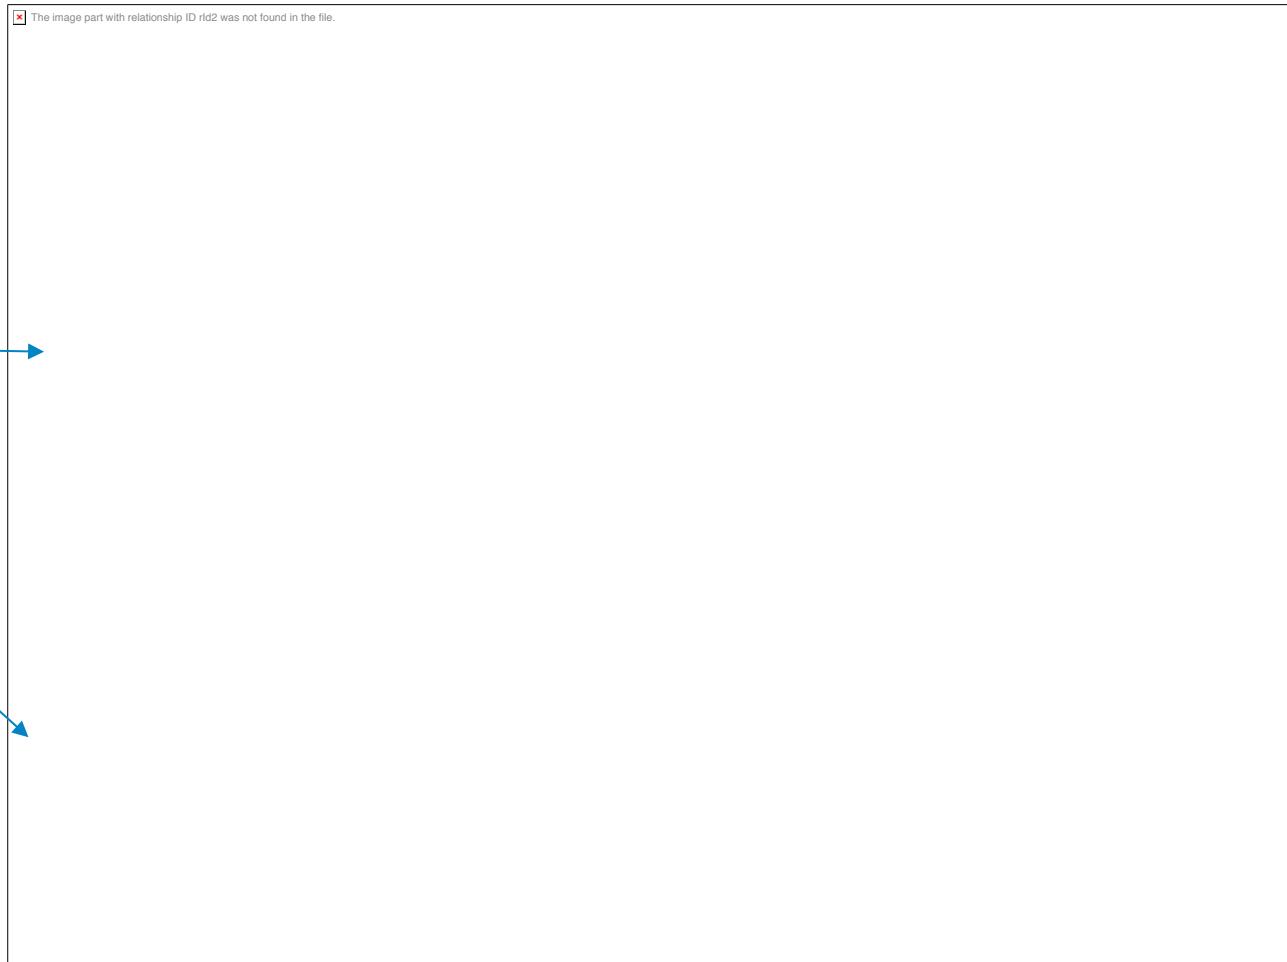
Like any list, this is an overview and

- Incomplete, e.g.
  - Business-APIs-as-a-Service: Payments is an example, but there are others, e.g.
    - Order Management, Inventory, Shipping, ...
    - Calendar and events. Contact management.
    - Human Resources.
  - Information-as-a-Service
    - Demographics.
    - Credit history and information.
    - Market data, prices, etc.
- Overly simplistic, e.g.
  - IaaS is many types of compute, storage, networking, ...
  - There are many, many types of PaaS
    - Integration-Platform-as-a-Service
    - More classical app/container, e.g. Cloud Foundry
    - Force.com and high-level component assembly.

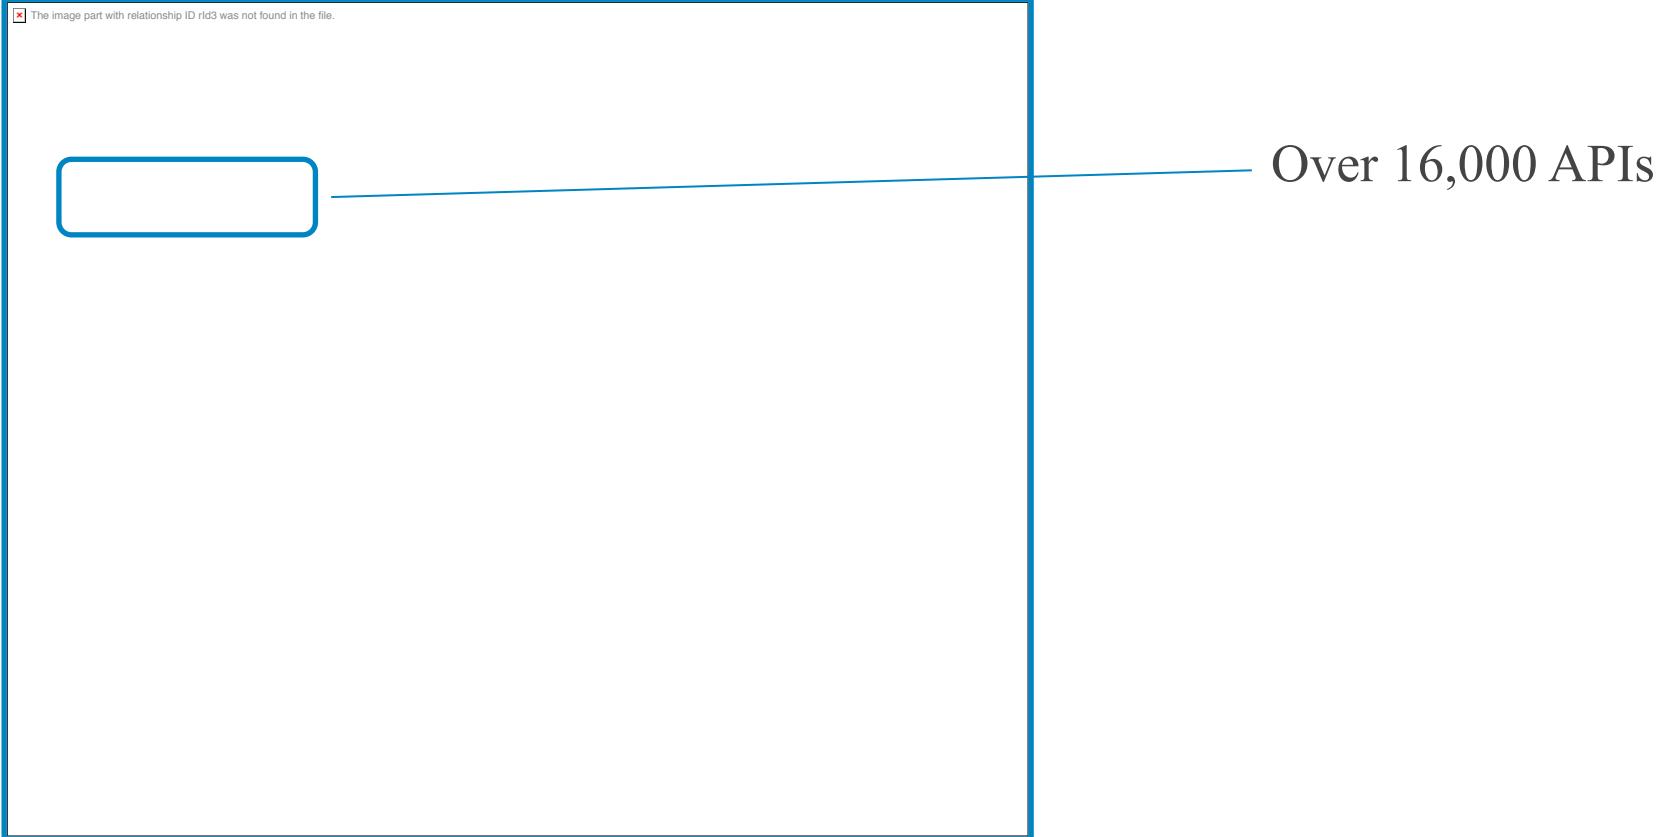
# Cloudscape

Our applications will

- Call APIs
- Use and run in PaaS



# www.programmableweb.com (Example)



# Microservices

# Microservices

(<https://martinfowler.com/articles/microservices.html>)

 The image part with relationship ID rid2 was not found in the file.

# Web Application Basic Concepts

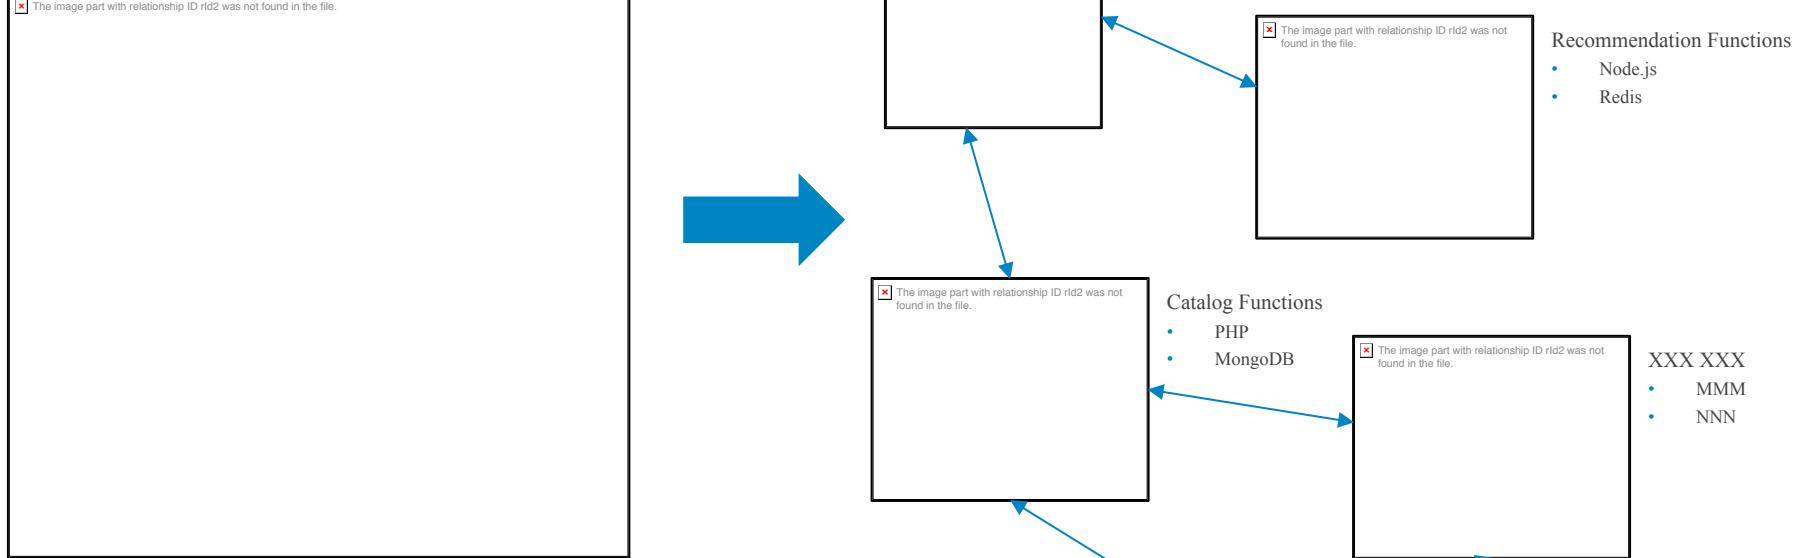
 The image part with relationship ID rId2 was not found in the file.

# “Traditional” Web Application Architecture



The image part with relationship ID rId2 was not found in the file.

# Monolithic to Micro



## Motivations

- Enable best tools, languages, ... for function.
- Simplifies change management and evolution.
- Better alignment of apps with business functions.
- Reuse of code and internet services.

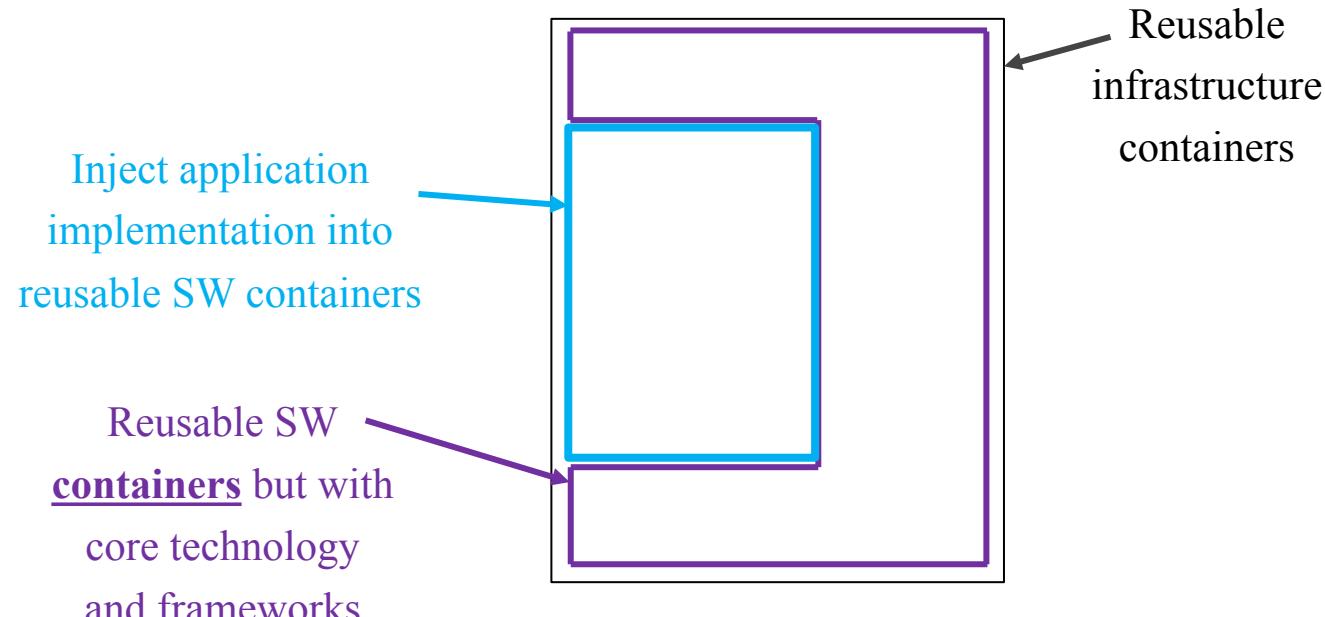
# Micro-services Characteristics

- Componentization via Services
- Organized around Business Capabilities
- Products not Projects
- Smart endpoints and dumb pipes
- Decentralized Governance
- Decentralized Data Management
- Infrastructure Automation
- Design for failure
- Evolutionary Design

# SOA vs Microservices

 The image part with relationship ID rId2 was not found in the file.

<http://usblogs.pwc.com/emerging-technology/agile-coding-in-enterprise-it-code-small-and-local/>



 The image part with relationship ID rId2 was not found in the file.

# Patterns

 The image part with relationship ID rid2 was not found in the file.

The image part with relationship ID rid2 was not found in the file.

## Patterns and practices are essential

- PaaS and advanced infrastructure eliminate the need to worry about a lot of things, but ...
- No infrastructure is smart enough to make a poorly designed application flexible, scalable, available, ...

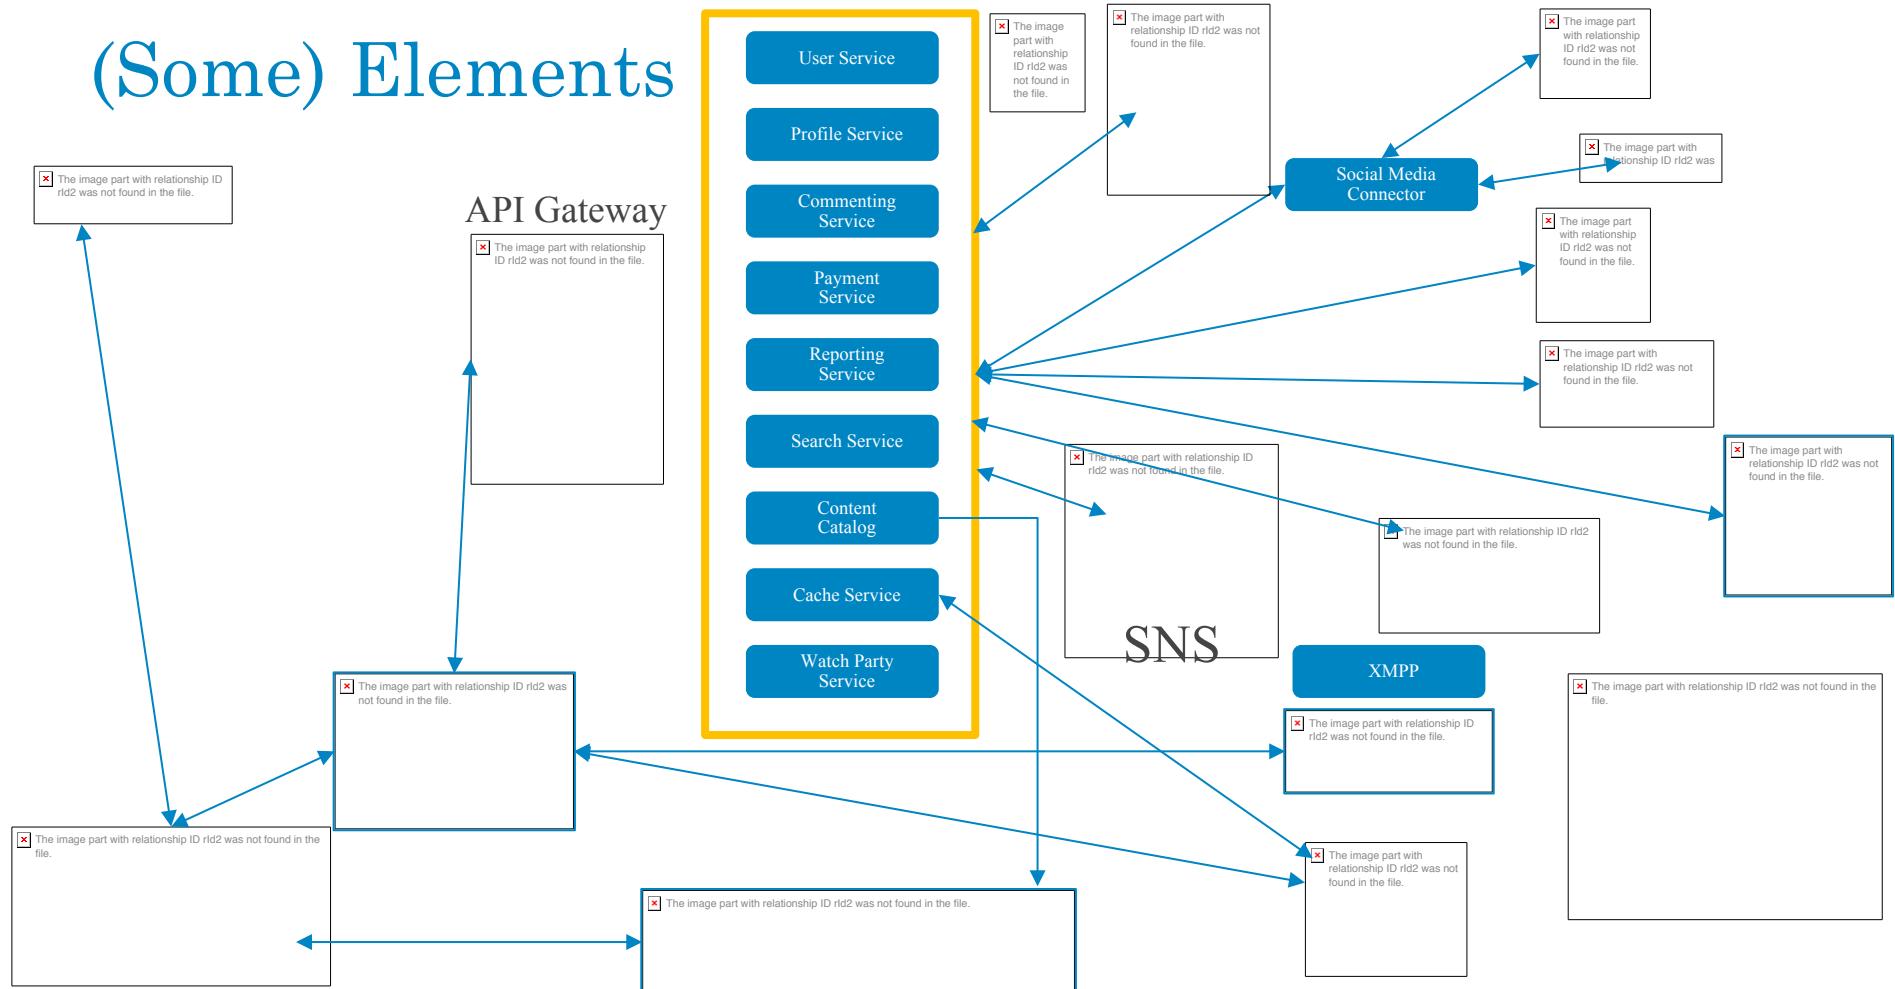
# Weinberg's Second Law

“If builders built buildings the way programmers wrote programs, then the first woodpecker that came along would destroy civilization.”

# An Example

# Seeka Demo

# (Some) Elements



# 0<sup>th</sup> Project

# 0<sup>th</sup> Project

- Teams
  - Form your teams (approx. 5 people)
  - Identify contact focal point.
  - Give your team a “cool” name.
- Signup (reuse) and Amazon Web Service Account
  - Free Tier should be fine.
  - Provide access to team members.
- Create an Elastic Beanstalk instance/application.
  - Use one of the sample application.
  - Will have to make more sophisticated starting next week;  
I will use Node JS with Express for Elastic Beanstalk examples.

# Project Structure

Server setup

The screenshot shows a Java IDE interface with the following components:

- Project Explorer:** Shows the project structure under "FirstMicroservice". Key folders include `bin`, `www`, `routes`, `services`, and `views`. A red arrow points from the text "Server setup" to the `www` folder.
- Code Editor:** Displays the `app.js` file with the following code:

```
#!/usr/bin/env node
/*
 * Module dependencies.
 */
var app = require('../app');
var debug = require('debug')('firstmicroservice:server');
var http = require('http');

/**
 * Get port from environment and store in Express.
 */
var port = normalizePort(process.env.PORT || '3001');
app.set('port', port);

/**
 * Create HTTP server.
 */
var server = http.createServer(app);

/**
 * Listen on provided port, on all network interfaces.
 */
server.listen(port);
server.on('error', onError);
server.on('listening', onListening);

/**
 * Normalize a port into a number, string, or false.
 */
function normalizePort(val) {
  var port = parseInt(val, 10);

  if (isNaN(port)) {
    return false;
  }

  if (port < 0 || port > 65535) {
    throw new Error(`Port ${val} is not a valid port.`);
  }

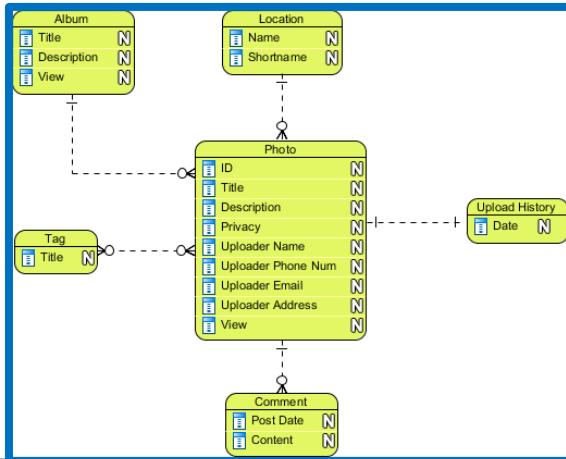
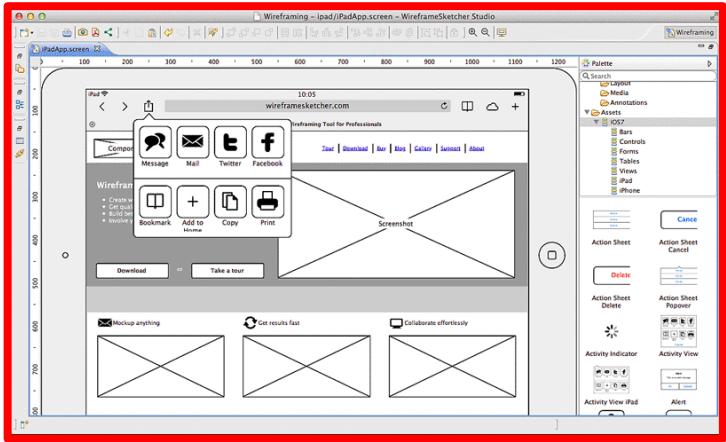
  return port;
}
```

- Run View:** Shows the command `/usr/local/bin/node /Users/donaldferguson/Dropbox/ColumbiaCourse/Courses/Fall2018/E6156/Projects/FirstMicroservice/bin/www` and the output of the application's logs.
- Bottom Status Bar:** Shows the status `In index.`

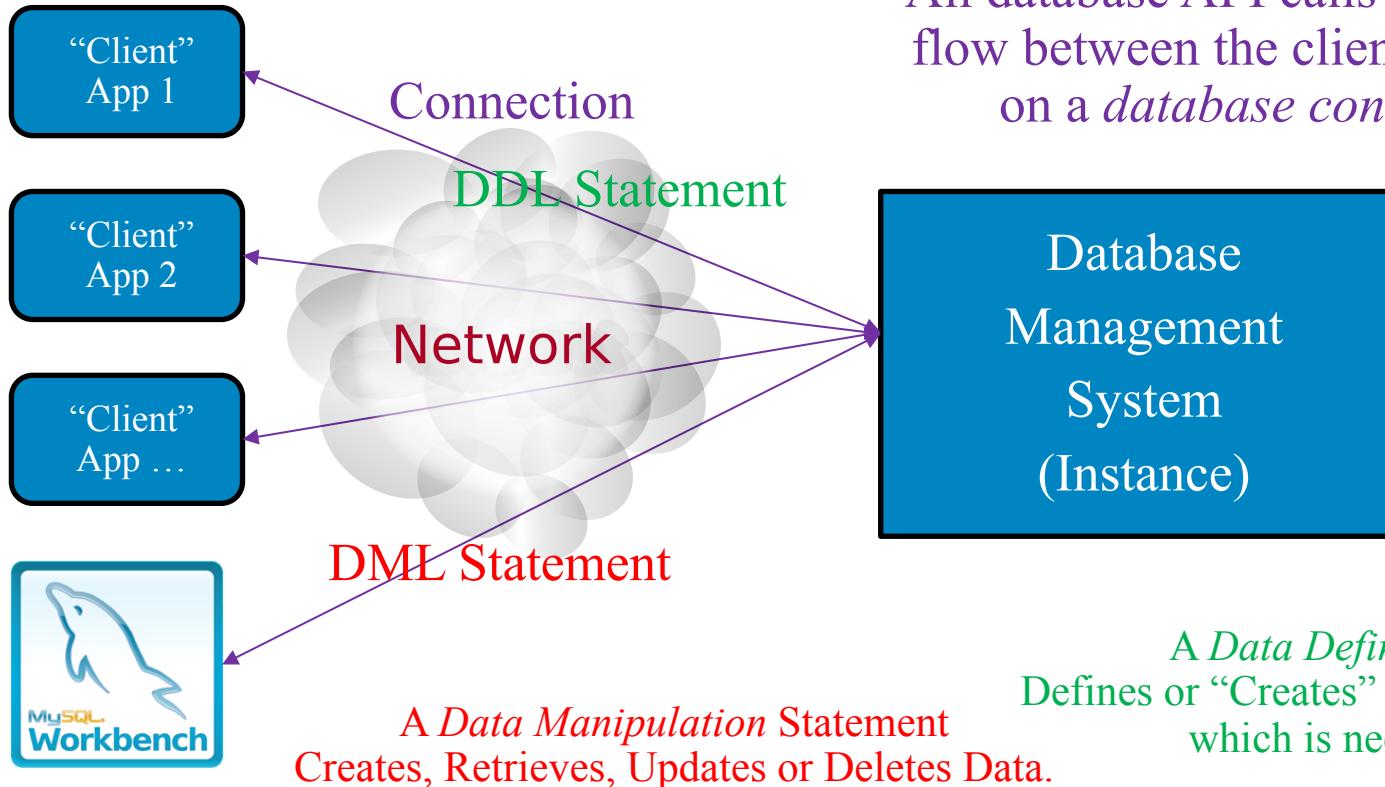
Application setup

# Application Design Methodologies (Simplistic, Somewhat Dated). (But, you get the idea).

- Start with User Interface
  - Roles and personas.
  - Wireframes for pages and interactions.
  - Page flow/transition diagrams.
  - Implement a model layer supporting UI
    - Functions
    - Data
- Start with Data Model
  - Entity types, properties.
  - Relationships/associations.
  - Constraints, integrity rules
  - Define reusable core operations.
  - Design role/task UIs that use core operations



# Database Connection

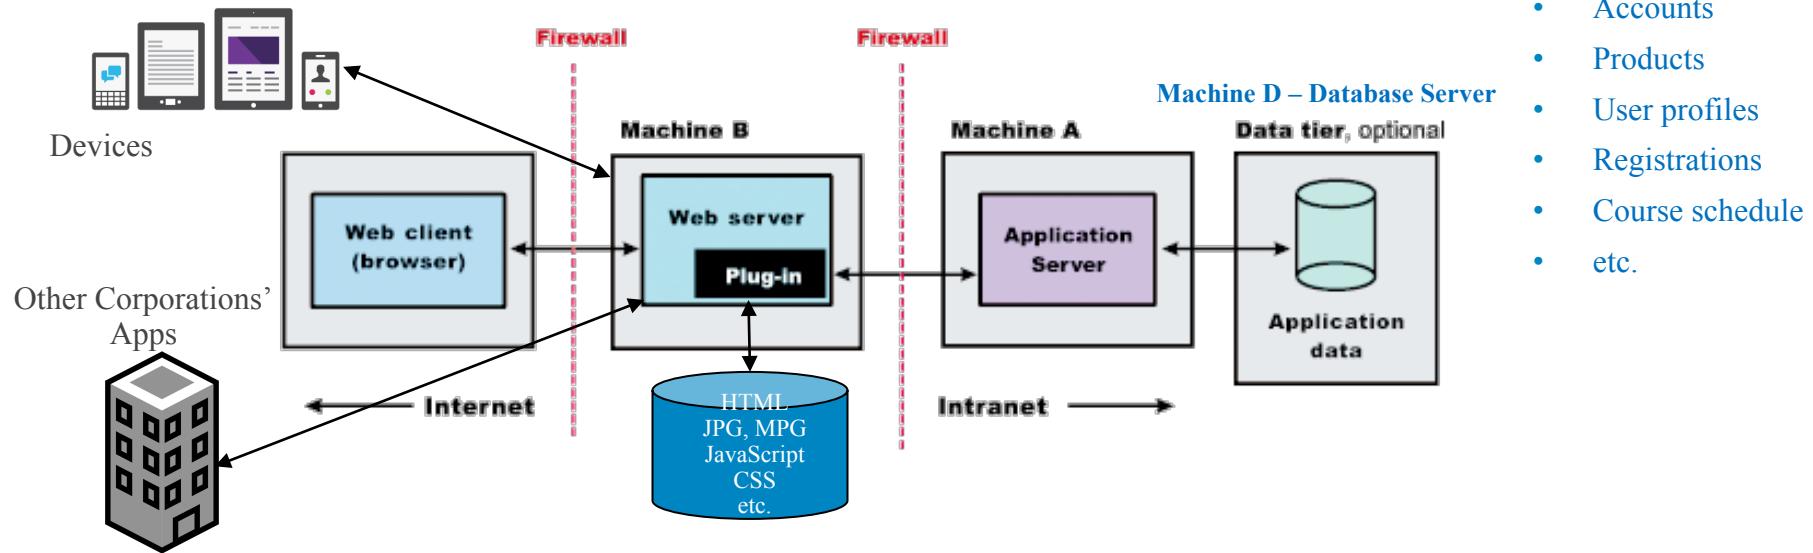


All database API calls (statements) flow between the client and server on a *database connection*.

Database  
Management  
System  
(Instance)

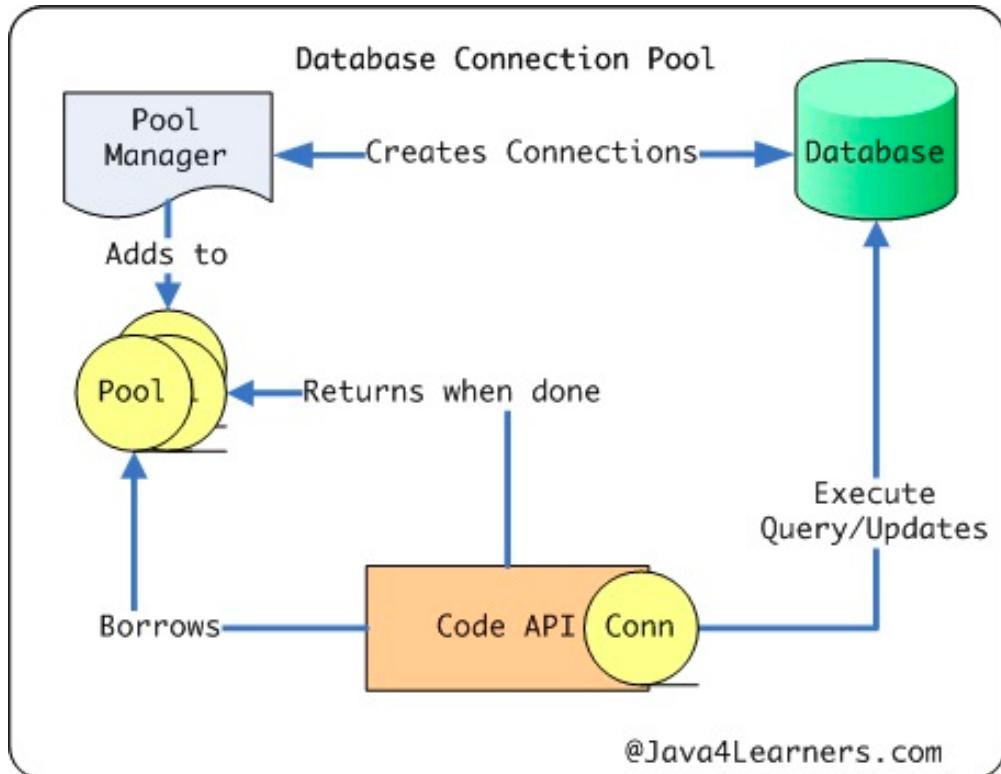
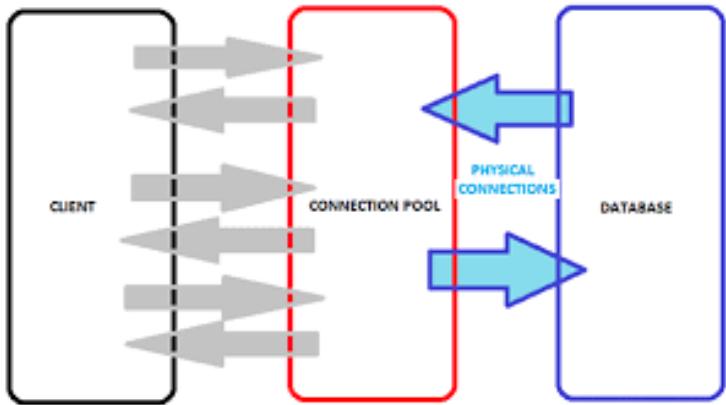
A *Data Definition Statement* Defines or “Creates” a definition/type of data, which is needed for DML.

# Web/Internet Application Structure



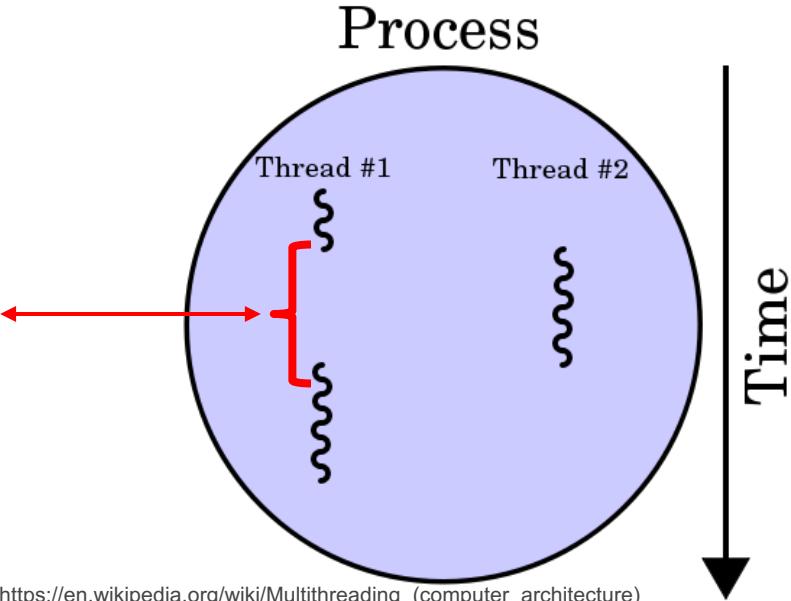
- There are 100s, 1000s, millions of simultaneous user sessions.
- There may be 100s, 1000s of servers at each tier to achieve
  - Scalability (throughput, response time)
  - Availability, fault tolerance.
- Servers may be spread over multiple data centers, public/private clouds, etc.

# Connection Pool



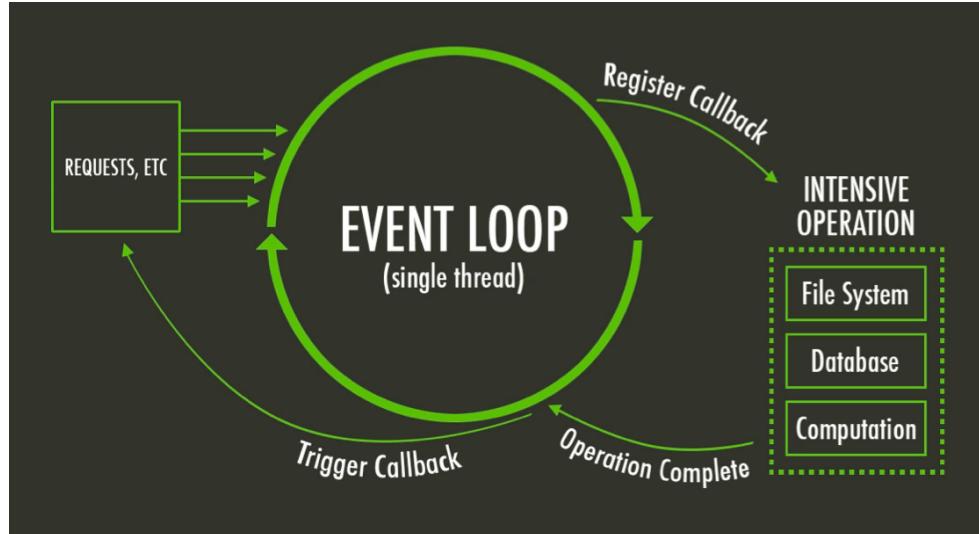
@Java4Learners.com

# Multithreading and Callbacks

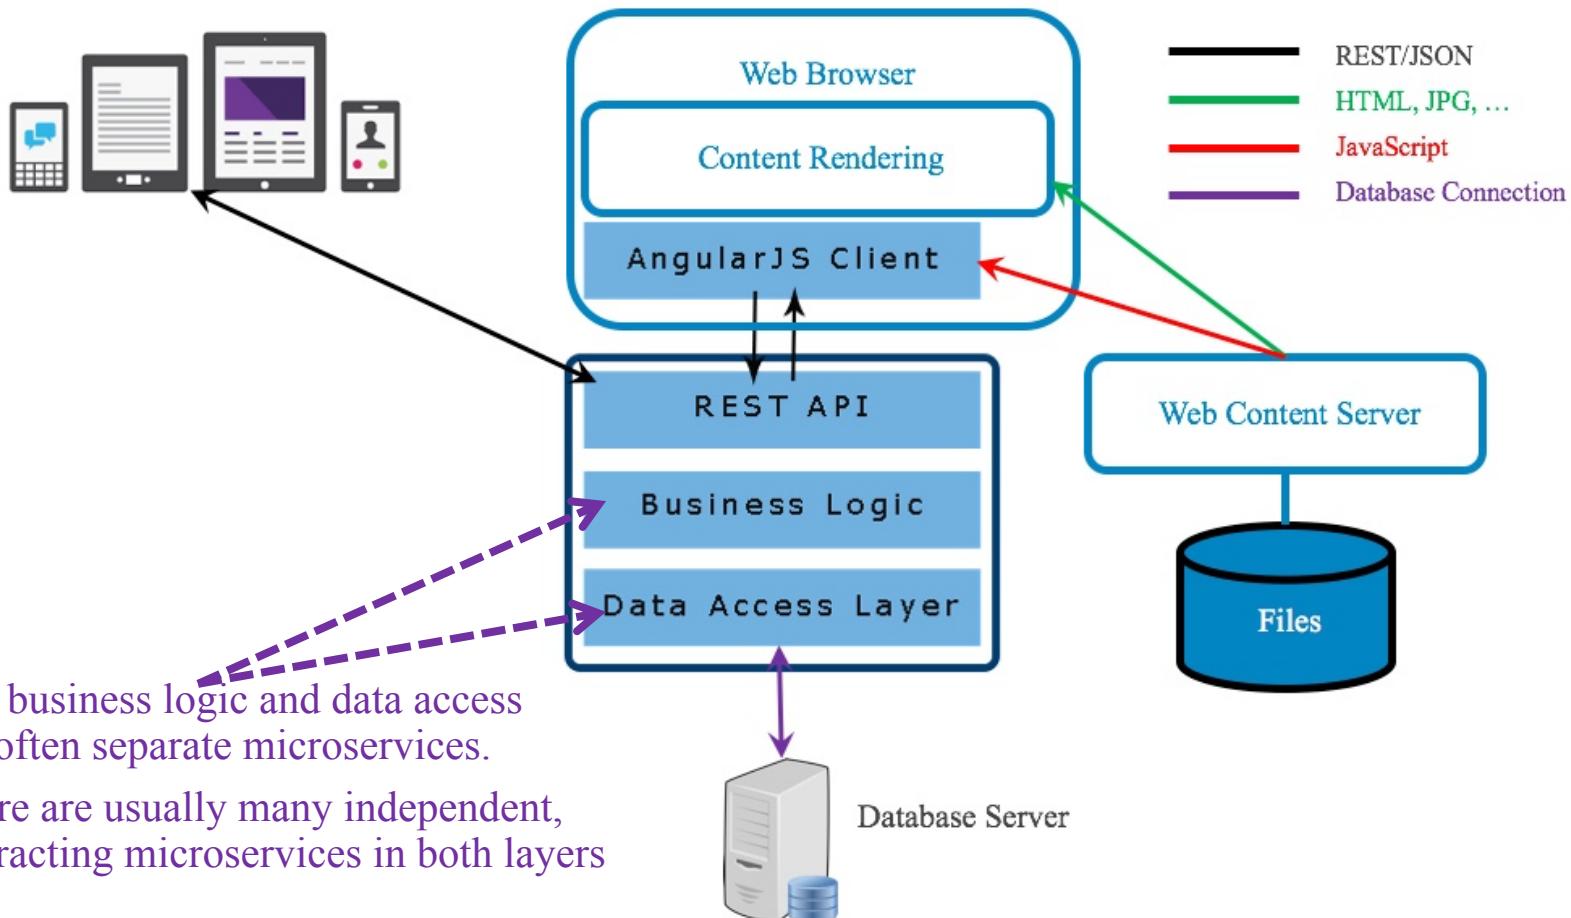


[https://en.wikipedia.org/wiki/Multithreading\\_\(computer\\_architecture\)](https://en.wikipedia.org/wiki/Multithreading_(computer_architecture))

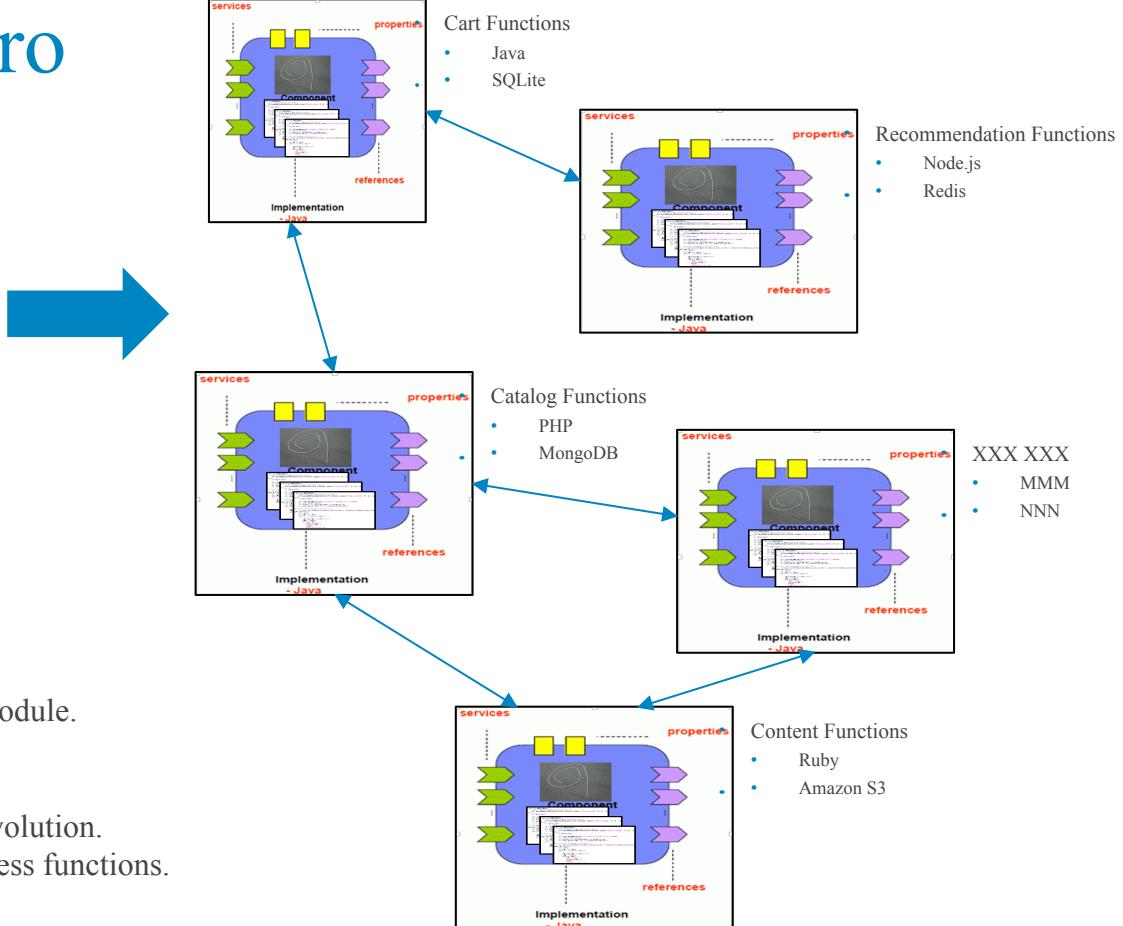
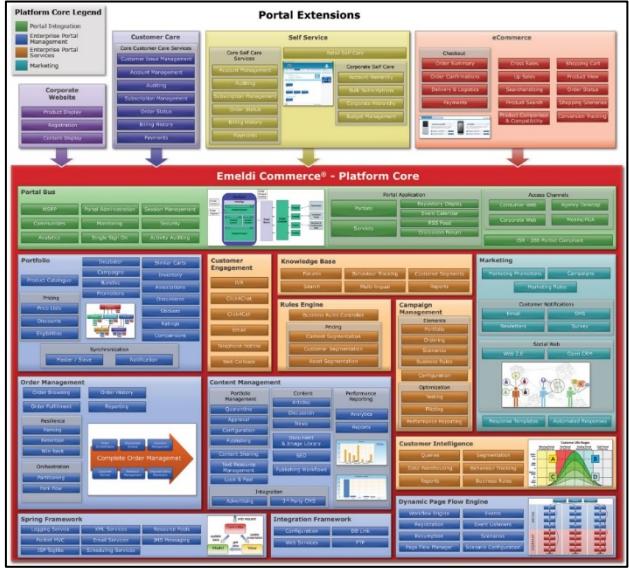
- Multithreading
  - Thread #1 “blocks” for an IO or API call.
  - Process switches to “ready” thread #2



- Callbacks
  - Incoming requests map to an event.
  - There is a single thread running application code associated with an event..
  - Long running calls get assigned to a separate worker thread, and register a callback.
  - Completion of an operation results in an event, which resumes the application via a callback.

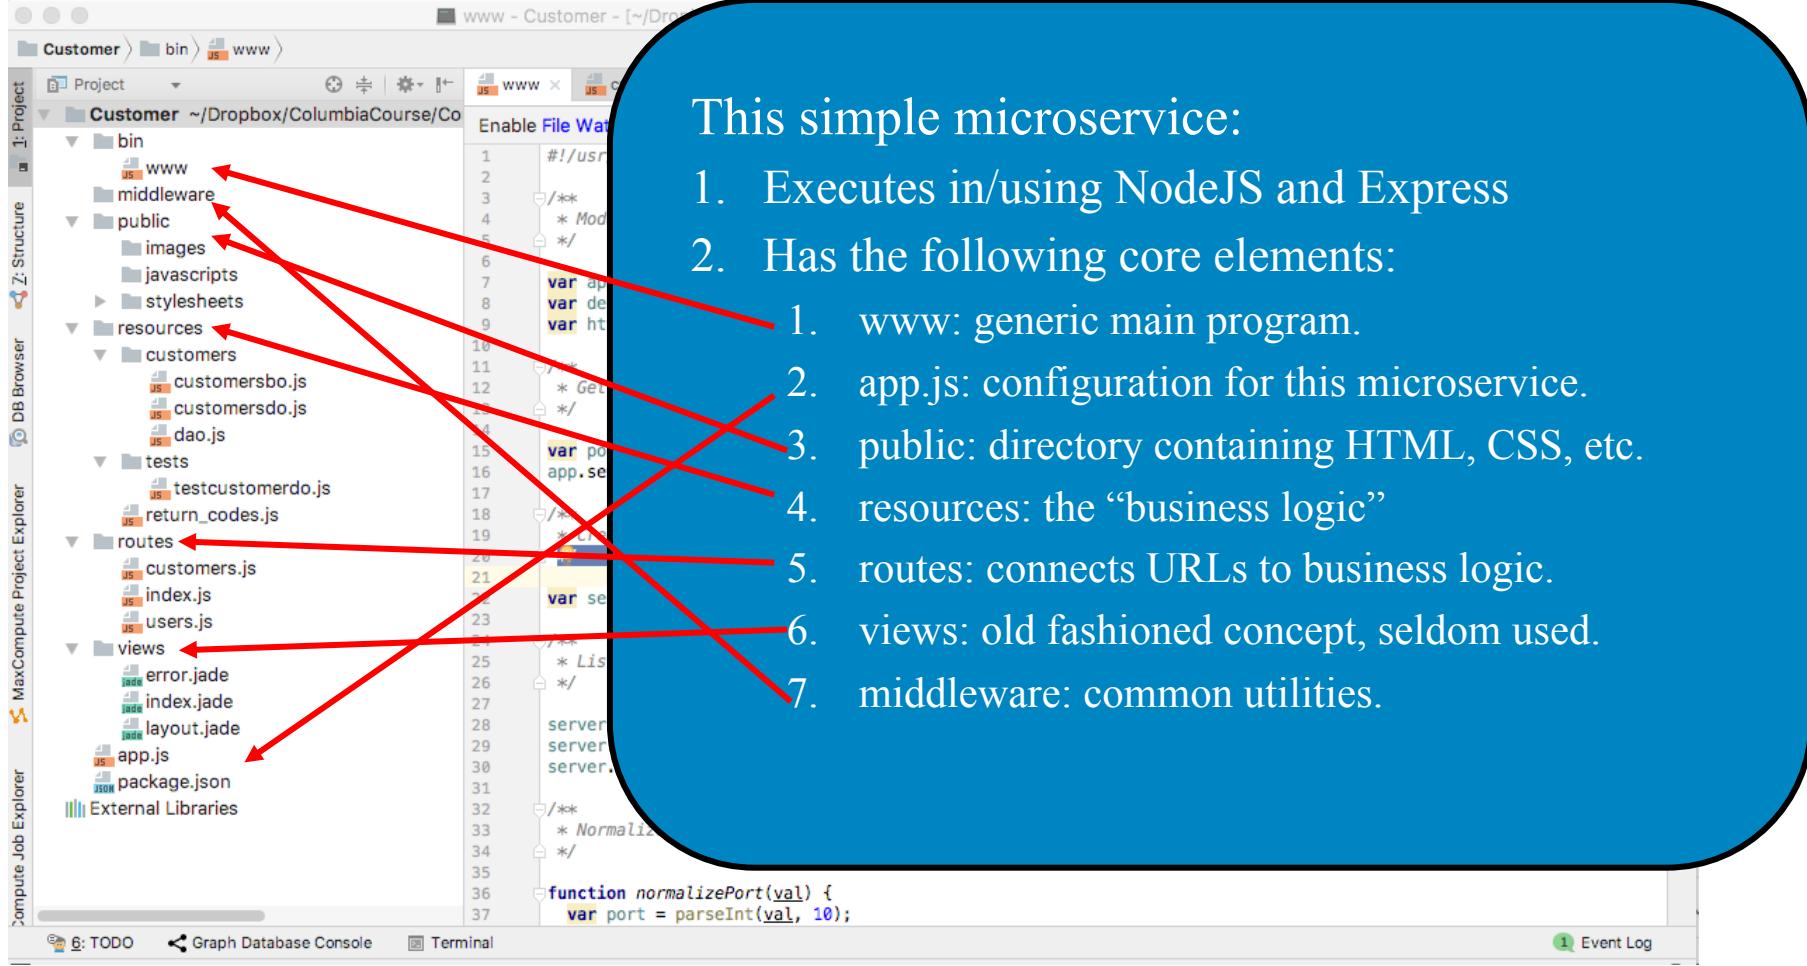


# Monolithic to Micro



## Motivations

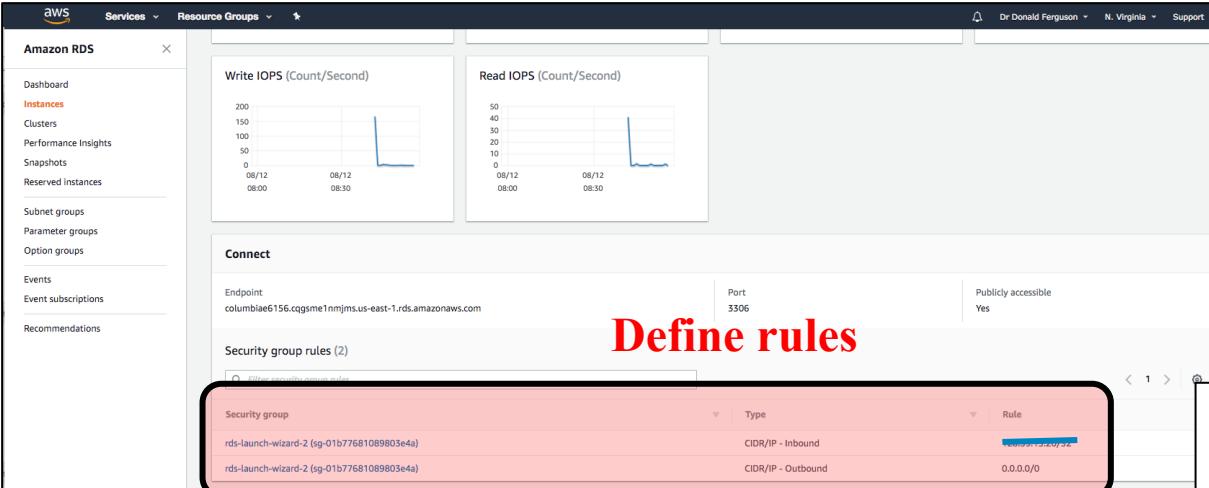
- Enable best tools, languages, ... for module.
  - Polyglot persistence
  - Polyglot programming
- Simplifies change management and evolution.
- Better alignment of “apps” with business functions.
- Reuse of code and internet services.



This simple microservice:

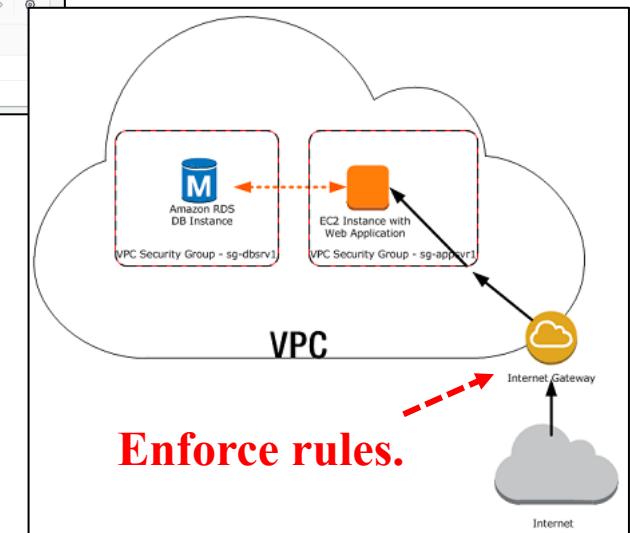
1. Executes in/using NodeJS and Express
2. Has the following core elements:
  1. www: generic main program.
  2. app.js: configuration for this microservice.
  3. public: directory containing HTML, CSS, etc.
  4. resources: the “business logic”
  5. routes: connects URLs to business logic.
  6. views: old fashioned concept, seldom used.
  7. middleware: common utilities.

# RDS/EC2 Security Groups



Define rules

- RDS (and EC2) instances are in a security group.
- Security groups has several functions, including
  - Preventing network access to group resources.
  - Unless the source address is on a whitelist.
- A gateway enforces these rules, including access from the public Internet.

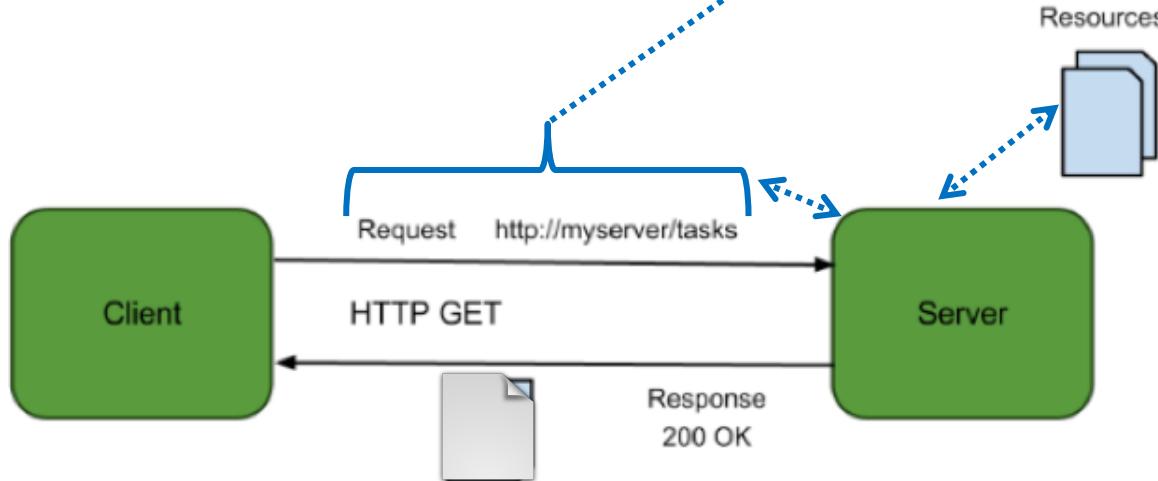


## Accept type in headers.

```
Accept: < MIME_type >/< MIME_subtype >
Accept: < MIME_type >/*
Accept: */*

// Multiple types, weighted with the quality value syntax:
Accept: text/html, application/xhtml+xml, application/xml;q=0.9, */*;q=0.8
```

- Relative URL identifies “resource” on the server.
- Server implementation maps abstract resource to tangible “thing,” file, DB row, ...

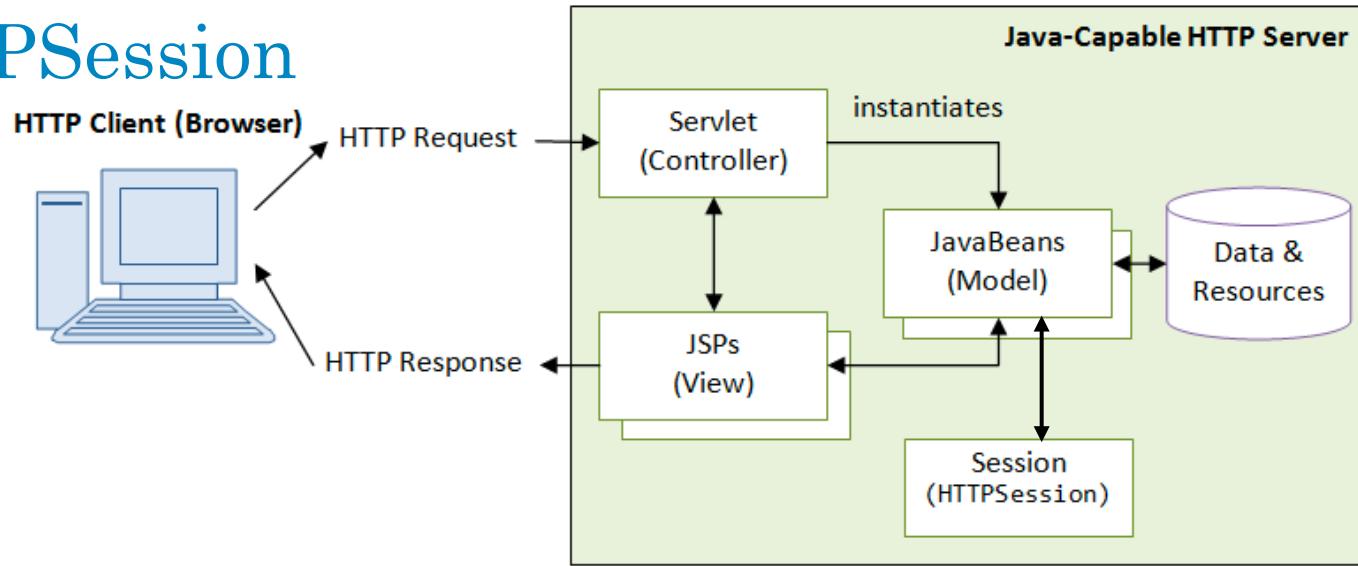


Client may be

- Browser
- Mobile device
- Microservice
- ... ...

Web media (content) type, e.g.  
• `text/html`  
• `application/json`

# HTTPSession



- Server can start a session based on login, first request from IP address, specific request, ...
- Session has an ID.
- Server can save data in the **HTTPSession** and access via ID.
- Server reads and updates **HTTPSession** on each request.
- The response to request N is a function of:
  - Request parameters.
  - Session information from previous requests.
  - Resource state



# Sample Facebook Operation

The screenshot shows the Postman application interface with the following details:

- Method:** GET
- URL:** graph.facebook.com/v3.1/me/friends?limit=2
- Headers (1):** Authorization (Value: Bearer EAAYEsPH7hkIBABloBMcxC1la5XnPfj75aO8vleTPbh...)
- Body:** JSON (Pretty, Raw, Preview) - Displays the API response.
- Test Results:** Status: 200 OK, Time: 167 ms

**API Response (Pretty JSON):**

```
1 {  
2   "data": [  
3     {  
4       "name": "Dw [REDACTED]",  
5       "id": "9496 [REDACTED]"  
6     },  
7     {  
8       "name": "Ch [REDACTED]",  
9       "id": "10102 [REDACTED]"  
10    }  
11  ],  
12  "paging": {  
13    "cursors": {  
14      "before": "QVFIUkYycFI3RH1elR6VHFwb0ZAFRUtaUUlGMGZAFYkNLWNNhb0V1MW40ZAkgrwaUNyblLUnNoMkFUaEnrQ25PSFpnWWzD",  
15      "after": "QVFIUuNzF10ERLN1VVtZAUNFVRNHAzdHN0MTZA1SU42MkpsR1V4dU5PR0VQM0g0SEl3cGhmc2t5MjIxS1NJTTThhVjYZD"  
16    },  
17    "next": "https://graph.facebook.com/v3.1/10153594875923693/friends?pretty=1&limit=2&after=QVFIUuNzF10ERLN1VVtZAUNFVRNHAzdHN0MTZA1SU42Mkp:  
18  },  
19  "summary": {  
20    "total_count": 783  
21  }  
22 }
```

# Cacheability

There are caches everywhere:

- Browser client.
- Enterprise perimeter.
- Content delivery network (CDN)
- ISP nodes.

